

6.S897/HST.S56 Problem Set 4

Released: Tues Mar 12

Due: Tues Mar 19 by 11:59pm through Gradescope

Instructions

You will find a *psets/pset4* directory in the class Github. Note this is different from prior problem sets. This repo contains the write-up, starter notebook, and the code used to extract the data. The data is available on [Physionet](#) after logging in through “Authorized users may enter here.”. Specifically you will need **demo.csv** and **vitals.csv** from Physionet and **train_ids.pk**, **valid_ids.pk**, and **test_ids.pk** from the Github. Please let us know ASAP if you have any issues accessing the data.

In the body of this exercise, [blue text](#) describes what specifically you should be submitting. If you have any questions about what the wording of the questions mean, ask us on Piazza!

For your final assessment, you should a [standalone report of your analyses with relevant plots, interpretation of findings, and recommendations named \$\\${mit_user}.pdf\$](#) (e.g. *iychen.pdf*) and any code used for the write-up.

In your write-up, please make your answers as easy to identify as possible (e.g. highlight with colors, label questions with numbers, etc). The faster we can identify your answers, the faster we can grade 70 submissions! :)

Part 1: Physiological signals can have signal [10 points]

We have provided vitals information from the MIMIC-III dataset for adult patients. Data was gathered using **make_vitals.py** which is included in the Github for those curious. In particular, we have provided vitals for patients from the first 24 hours for the following vitals

- Heart rate in bpm
- Respiratory rate in breaths / minute
- Mean arterial pressure
- Pulse oximetry
- Systolic blood pressure in mmHg
- Diastolic blood pressure in mmHg
- Glucose reading
- Temperature

We have aggregated each vital signal into the min/max/mean of that hour. Because not all vital signs are measured at the same rate, missing values are designated as NaN. If no vital signs were measured in that hour, the patient ICU stay will have no row for that hour.

1.1 How many rows are in `vitals_24h.csv`? For each of the vital types, how many vital signs are there per patient ICU stay on average? Plot the distribution of values across patient ICU stays for each vital sign using histograms with 40 bins.

1.2 We want to model patients with ICU stays where we have hourly readings for the first 24 hours. Patients with ICU stays where measures aren't recorded have been marked as null.

- How many patient ICU stays are included in this dataset in any form?
- How many patients ICU stays show up with 24 hours of measurements, even if some of the hours have some missing values?
- How many patients ICU stays have 24 hours measurements with all present measurements?

1.3 In order to increase the number of include patient ICU stays, we want to keep vital readings that appear >80% of the time over all rows in `vitals_24h.csv`. Which vitals would we keep?

1.4 For the rest of the problem set, we will keep vitals that appear >80% of the time across all rows, but that is only one way to handle missing data. Besides throwing out rows that are incomplete, what are three other ways we can impute missing data for multivariate time series modeling?

Part 2: Long short term memory [20 points]

2.1 We are interested in using this time series information to model patient care. Keep vitals that appear in >80% the rows. (As a sanity check, this should include 18 columns.) [How many patient ICU stays are in the dataset?](#)

We want to analyze patient ICU stays that last longer than 48 hours by predicting from the first 24 hours, but `vitals_24h.csv` includes vitals from all patients in the MIMIC dataset. Join with `demo.csv`, which only includes patients ICU stays that last longer than 48 hours. [How many patients are left in the dataset now?](#)

Create a data tensor of shape: `[n_samples, n_timesteps, n_features]` for data labeled train, valid, and test. Data should be scaled using `MinMaxScaler` from `sklearn` using `feature_range=(0,1)` across the entire dataset before being separated into train, valid, and test.

You may find it helpful to create column names of the form: `t00_hearttrate_min`, `t00_hearttrate_max`, ..., `t23_hearttrate_max`, ...

We often find that sex and age are predictive features. Encode sex as a "is_female" feature. Include age and sex in the tensor: `t00_is_female`, `t00_age`, ..., `t23_is_female`, `t23_age`. Note that `t00_age`, ..., `t23_age` will all be the same and similar with the sex features. These features can be found in `demo.csv`.

We have given you the `icustay_ids` corresponding to the train, valid, and test sets in **`train_ids.pk`**, **`valid_ids.pk`**, and **`test_ids.pk`**, which you can load using the pickle package in Python. If you are having trouble loading these IDs, or they do not match your dataframe, make a note in your write-up and describe your process of randomly selecting training/validation/testing data in a 60/20/20 split.

2.2 Build an LSTM using Keras for multivariate time series classification, predicting hospital mortality or `mort_hosp`. Use an LSTM with:

- 20 units
- a dense layer with sigmoid activation
- binary cross-entropy loss
- ADAM optimizer
- Dropout of 0.2
- 100 epochs
- Batch size of 128

You may find these guides helpful:

- [Multivariate time series forecasting for regression](#)
- [Sequence classification using LSTMs](#)

What is the train loss, accuracy, and AUC? Experiment with Dropout = [0.0, 0.2, 0.5, 0.8] by comparing the AUC on the validation data. What Dropout value yields the best AUC on the validation data? Using this best Dropout value, what is the corresponding test AUC?

2.3 Select two patient ICU stays for whom the LSTM predicts the patient ICU stay will end in patient mortality and two patient ICU stays for whom the LSTM predicts the patient ICU stay will NOT end in patient mortality. [Plot any vital features that are particularly meaningful over time and describe any noticeable differences.](#)

Part 3: Compare to linear baseline [10 points]

Whenever we use more sophisticated deep learning models, we want to compare against a simpler baseline to make sure the computational effort is worth it.

3.1 [What is Logistic Regression with just the raw values from 2.2?](#) You should use a matrix of shape `[n_shapes, n_timesteps * n_features]`, which is an unrolled version of the data tensor from earlier. You can keep age and sex in their time-multiplied form or only include them once.

3.2 Hyperparameter tune on `C=[0.01, 0.1, 0.5, 1.0, 5.0]` and `penalty=['l1', 'l2']` by comparing the AUC on the validation data. [What are the best hyperparameters on the validation data? What is the test AUC using these best hyperparameters? How does this compare to the results from 2.2?](#)

Part 4: Not graded by please answer [0 points]

4.1 How many hours did you spend on this problem set?

4.2 This past week (Tues March 12 - Tues March 19), how many hours did you spend (attending lectures, reading papers, problem sets) on the class?